

Spectral Admissibility and Non-Abelian Mode Resilience in the Q_8 Cayley Graph under Bounded Flux

Working note – Cosmochrony programme

March 7, 2026

Abstract

We establish the first analytically tractable instance of the spectral admissibility programme outlined in the Cosmochrony framework. On the undirected Cayley graph of the quaternion group Q_8 , with generators $S = \{\pm i, \pm j, \pm k\}$, the local bounded-flux constraint $|\partial_t \chi_v| \leq c_\chi$ induces a hierarchy of maximal admissible effective amplitudes across irreducible spectral sectors.

In the linearised dispersion regime, the non-abelian irreducible sector admits a maximal effective amplitude larger than the nontrivial abelian sectors by a factor of $\sqrt{4/3} \approx 1.155$. Resolving explicitly the basis dependence in the degenerate ρ_4 eigenspace, we show that this advantage holds for the physically relevant effective amplitude $A_n = a_n \|\psi_n\|_\infty$, but not for the raw amplitude a_n .

We then perform the weakly non-linear modal reduction of the discrete Dirac–Born–Infeld (DBI) Lagrangian and derive the first amplitude-dependent correction to the dispersion relation, governed by the quartic shape factor $\kappa_n = \sum_v \psi_n(v)^4$. When expressed in terms of the effective amplitude, this correction depends only on the Laplacian eigenvalue λ_n , leaving the spectral hierarchy unchanged.

Finally, solving the one-mode DBI dynamics exactly on the Q_8 Cayley graph, we show that the admissibility envelope

$$A_n^{\max} = \frac{c_\chi}{\sqrt{\lambda_n}}$$

coincides with the exact saturation boundary of the DBI dynamics. The spectral hierarchy therefore remains structurally stable beyond the perturbative regime.

This provides a minimal explicit demonstration that bounded substrate capacity can favour non-abelian relational modes purely through spectral admissibility, without invoking additional dynamical selection principles.

More broadly, it suggests that non-abelian symmetry sectors may emerge as the dynamically most resilient components of the relational spectrum under bounded-flux constraints.

1 Setting and Motivation

The Cosmochrony framework identifies the physical spectrum with the residue of two successive constraints acting on the relational operator $\mathcal{L}_\chi$:

1. **Topological discretisation:** the fibre structure of the non-injective projection Π organises modes into discrete winding sectors.
2. **Bounded-flux filtering:** the saturation bound $|\partial_t \chi| \leq c_\chi$ eliminates modes whose local evolution would exceed the substrate’s relational capacity.

The goal of this note is to make step 2 explicit and calculable on the simplest non-trivial finite graph encoding a quaternionic algebraic structure, as a minimal discrete proxy for $SU(2)$ -type sectors.

2 The Q_8 Cayley Graph

Let $Q_8 = \{\pm 1, \pm i, \pm j, \pm k\}$ denote the quaternion group of order 8. We define the undirected Cayley graph $\Gamma(Q_8, S)$ with generating set

$$S = \{\pm i, \pm j, \pm k\}. \quad (1)$$

Since S is closed under inversion ($s^{-1} \in S$ for all $s \in S$) and does not contain the identity, $\Gamma(Q_8, S)$ is a well-defined, connected, undirected, 6-regular graph on 8 vertices.

Remark 1. *The choice $S = \{i, j, k\}$ (without inverses) would yield a directed graph with an asymmetric adjacency operator. For a self-adjoint Laplacian, the full symmetric set $S = \{\pm i, \pm j, \pm k\}$ is required. The resulting graph is 6-regular, not 3-regular.*

2.1 Adjacency Operator and Graph Laplacian

The adjacency operator A of $\Gamma(Q_8, S)$ acts on functions $f: Q_8 \rightarrow \mathbb{C}$ by

$$(Af)(g) = \sum_{s \in S} f(gs) = \sum_{s \in \{\pm i, \pm j, \pm k\}} f(gs). \quad (2)$$

The combinatorial graph Laplacian is

$$\mathcal{L} = dI - A, \quad d = |S| = 6, \quad (3)$$

where I is the identity on $\ell^2(Q_8)$. The eigenvalues of $\mathcal{L}$ are $\lambda = 6 - \mu$, where μ ranges over the eigenvalues of A .

3 Spectral Decomposition via Representation Theory

Since Q_8 is a finite group, $\ell^2(Q_8) \cong \mathbb{C}[Q_8]$ decomposes as a Q_8 -module into irreducible representations. By the Peter–Weyl theorem for finite groups:

$$\mathbb{C}[Q_8] \cong \bigoplus_{\rho \in \widehat{Q_8}} V_\rho^{\oplus \dim \rho}, \quad (4)$$

where the sum runs over the five irreducible representations of Q_8 .

3.1 Irreducible Representations of Q_8

Q_8 has five irreducible representations:

Label	Dimension	$\rho(i)$	$\rho(j)$	$\rho(k)$
ρ_0 (trivial)	1	1	1	1
ρ_1	1	1	-1	-1
ρ_2	1	-1	1	-1
ρ_3	1	-1	-1	1
ρ_4 (quaternionic)	2	$\begin{pmatrix} i & 0 \\ 0 & -i \end{pmatrix}$	$\begin{pmatrix} 0 & 1 \\ -1 & 0 \end{pmatrix}$	$\begin{pmatrix} 0 & i \\ i & 0 \end{pmatrix}$

The dimension check: $1^2 + 1^2 + 1^2 + 1^2 + 2^2 = 8 = |Q_8|$. ✓

3.2 Eigenvalues of the Adjacency Operator

For any irreducible representation ρ , the adjacency operator acts as the scalar $\mu_\rho \cdot I_{V_\rho}$ on V_ρ . Since S is a union of conjugacy classes of Q_8 , the corresponding element of the group algebra is central, and Schur's lemma gives

$$\mu_\rho = \frac{1}{\dim \rho} \sum_{s \in S} \chi_\rho(s), \quad (5)$$

where χ_ρ is the character of ρ .

1D sectors. For the trivial representation ρ_0 : $\chi_{\rho_0}(s) = 1$ for all $s \in S$, so $\sum_{s \in S} \chi_{\rho_0}(s) = 6$ and

$$\mu_{\rho_0} = \frac{1}{1} \cdot 6, \quad \lambda_{\rho_0} = 6 - 6 = 0. \quad (6)$$

For ρ_1 : $\chi_{\rho_1}(\pm i) = 1$ and $\chi_{\rho_1}(\pm j) = \chi_{\rho_1}(\pm k) = -1$, so $\sum_{s \in S} \chi_{\rho_1}(s) = 2(1) + 4(-1) = -2$, giving $\mu_{\rho_1} = -2$ and $\lambda_{\rho_1} = 8$. By the symmetry of Q_8 under permutations of $\{i, j, k\}$, $\mu_{\rho_2} = \mu_{\rho_3} = -2$ and $\lambda_{\rho_2} = \lambda_{\rho_3} = 8$.

2D sector. For the quaternionic representation ρ_4 , we compute the character at each generator. The matrices $\rho_4(i)$, $\rho_4(-i)$, $\rho_4(j)$, $\rho_4(-j)$, $\rho_4(k)$, $\rho_4(-k)$ each have trace zero: for instance, $\text{tr}(\rho_4(i)) = \text{tr}\begin{pmatrix} i & 0 \\ 0 & -i \end{pmatrix} = 0$, and $\text{tr}(\rho_4(-i)) = 0$; similarly for j and k . Thus $\sum_{s \in S} \chi_{\rho_4}(s) = 0$ and

$$\mu_{\rho_4} = \frac{1}{2} \cdot 0 = 0, \quad \lambda_{\rho_4} = 6 - 0 = 6. \quad (7)$$

3.3 Summary of the Spectrum

Sector	λ_n	Multiplicity	Type
ρ_0	0	1	Trivial (zero mode)
ρ_1, ρ_2, ρ_3	8	3	Nontrivial abelian
ρ_4	6	4	Non-abelian (quaternionic)

Total dimension: $1 + 3 + 4 = 8$. ✓

4 Bounded-Flux Admissibility Condition

4.1 Setup

Let the substrate field $\chi: Q_8 \times \mathbb{R} \rightarrow \mathbb{R}$ evolve on the graph, subject to the local bounded-flux constraint

$$|\partial_t \chi_v(t)| \leq c_\chi \quad \forall v \in Q_8, \forall t. \quad (8)$$

Consider a single-mode excitation

$$\chi_v(t) = a_n \psi_n(v) \cos(\omega_n t), \quad (9)$$

where ψ_n is an eigenmode of $\mathcal{L}$ with eigenvalue λ_n , normalised in $\ell^2(Q_8)$, and $a_n > 0$ is the mode amplitude. The time derivative at vertex v is

$$\partial_t \chi_v(t) = -a_n \omega_n \psi_n(v) \sin(\omega_n t). \quad (10)$$

The constraint (8) requires

$$|\omega_n| \cdot a_n \cdot |\psi_n(v)| \leq c_\chi \quad \forall v \in Q_8. \quad (11)$$

The binding constraint is at the vertex where $|\psi_n(v)|$ is maximal:

$$|\omega_n| \cdot a_n \cdot \|\psi_n\|_\infty \leq c_\chi. \quad (12)$$

4.2 Effective Amplitude and Admissibility Envelope

Definition 1. The effective amplitude of a mode (ψ_n, a_n) is

$$A_n := a_n \|\psi_n\|_\infty. \quad (13)$$

A mode is dynamically admissible (in the linearised regime) if

$$|\omega_n| \cdot A_n \leq c_\chi. \quad (14)$$

In the linearised dispersion approximation, $\omega_n^2 \approx \lambda_n$, so the maximal admissible effective amplitude is

$$A_n^{\max} = \frac{c_\chi}{\sqrt{\lambda_n}}. \quad (15)$$

This defines the *admissibility envelope* in the (λ_n, A_n) plane.

Remark 2. The linearised dispersion $\omega_n^2 = \lambda_n$ is valid for $A_n \ll c_\chi/\sqrt{\lambda_n}$. The non-linear (Born-Infeld) correction modifies this relation near saturation and requires a modal reduction of the DBI action; it is deferred to a separate calculation. The envelope (15) should therefore be understood as the admissibility boundary in the weakly non-linear regime.

4.3 Explicit Eigenmodes in the ρ_4 Sector

The non-abelian sector ρ_4 has multiplicity 4 in $\ell^2(Q_8)$, so the quantity $\|\psi_n\|_\infty$ is not determined until an explicit basis of the eigenspace is chosen. A natural orthonormal basis is provided by the matrix coefficients of the quaternionic irreducible representation.

For a finite group G and an irreducible representation ρ , the standard Peter–Weyl basis of the corresponding isotypic component is given by

$$\psi_{uv}(g) = \sqrt{\frac{\dim \rho}{|G|}} \rho(g)_{uv}, \quad u, v \in \{1, \dots, \dim \rho\}. \quad (16)$$

For Q_8 and ρ_4 , this becomes

$$\psi_{uv}(g) = \frac{1}{2} \rho_4(g)_{uv}, \quad u, v \in \{1, 2\}, \quad (17)$$

since $\dim \rho_4 = 2$ and $|Q_8| = 8$.

Using the explicit matrices

$$\rho_4(1) = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}, \quad \rho_4(-1) = \begin{pmatrix} -1 & 0 \\ 0 & -1 \end{pmatrix}, \quad (18)$$

$$\rho_4(i) = \begin{pmatrix} i & 0 \\ 0 & -i \end{pmatrix}, \quad \rho_4(j) = \begin{pmatrix} 0 & 1 \\ -1 & 0 \end{pmatrix}, \quad \rho_4(k) = \begin{pmatrix} 0 & i \\ i & 0 \end{pmatrix}, \quad (19)$$

one finds that the diagonal modes ψ_{11}, ψ_{22} are supported on $\{\pm 1, \pm i\}$, while the off-diagonal modes ψ_{12}, ψ_{21} are supported on $\{\pm j, \pm k\}$. In all four cases,

$$|\psi_{uv}(g)| \in \left\{0, \frac{1}{2}\right\}, \quad (20)$$

hence

$$\|\psi_{uv}\|_\infty = \frac{1}{2}, \quad u, v \in \{1, 2\}. \quad (21)$$

For the one-dimensional nontrivial abelian sectors ρ_1, ρ_2, ρ_3 , the normalised eigenmodes are the corresponding characters scaled by $|Q_8|^{-1/2}$. Since these characters take values in $\{\pm 1\}$, one has

$$\|\psi_{\rho_k}\|_\infty = \frac{1}{\sqrt{8}}, \quad k = 1, 2, 3. \quad (22)$$

4.4 Basis-Resolved Admissibility

Proposition 1 (Basis-resolved admissibility on Q_8). *In the canonical Peter–Weyl basis of $\ell^2(Q_8)$, the maximal admissible effective amplitudes satisfy*

$$A_{\rho_4}^{\max} = \frac{c_\chi}{\sqrt{6}} > \frac{c_\chi}{\sqrt{8}} = A_{\rho_k}^{\max}, \quad k = 1, 2, 3, \quad (23)$$

whereas the corresponding maximal raw amplitudes satisfy

$$a_{\rho_4}^{\max} = \frac{c_\chi}{\sqrt{6} \|\psi_{\rho_4}\|_\infty} = \frac{2}{\sqrt{6}} c_\chi \approx 0.816 c_\chi, \quad (24)$$

$$a_{\rho_k}^{\max} = \frac{c_\chi}{\sqrt{8} \|\psi_{\rho_k}\|_\infty} = c_\chi, \quad k = 1, 2, 3. \quad (25)$$

Thus the non-abelian sector is favoured in effective amplitude, but not in raw amplitude.

Proof. The effective-amplitude statement follows from $A_n^{\max} = c_\chi / \sqrt{\lambda_n}$ together with $\lambda_{\rho_4} = 6$ and $\lambda_{\rho_k} = 8$. The raw-amplitude statement follows from $a_n^{\max} = A_n^{\max} / \|\psi_n\|_\infty$, using $\|\psi_{\rho_4}\|_\infty = 1/2$ and $\|\psi_{\rho_k}\|_\infty = 1/\sqrt{8}$. $\square$

Remark 3 (Interpretation). *The bounded-flux constraint acts on the local quantity $|\omega_n| a_n |\psi_n(v)|$, and therefore selects modes through the effective amplitude $A_n = a_n \|\psi_n\|_\infty$, not through the raw amplitude a_n alone. The non-abelian advantage on Q_8 is therefore not a larger absolute amplitude capacity, but a larger admissible effective window per unit spectral weight. Equivalently, the ρ_4 sector is less constrained by bounded-flux filtering in the physically relevant local variable.*

4.5 Weakly Non-linear Modal DBI Reduction

To incorporate the first Born–Infeld correction, consider the discrete DBI Lagrangian

$$L = -c_\chi^2 \sum_v \sqrt{1 - \dot{\chi}_v^2 / c_\chi^2} - \frac{1}{2} \sum_{v,w} \chi_v \mathcal{L}_{vw} \chi_w. \quad (26)$$

For a single eigenmode,

$$\chi_v(t) = q(t) \psi_n(v), \quad (27)$$

with $\mathcal{L}\psi_n = \lambda_n \psi_n$ and $\|\psi_n\|_2 = 1$, this reduces to the exact one-mode Lagrangian

$$L_n[q] = -c_\chi^2 \sum_v \sqrt{1 - \psi_n(v)^2 \dot{q}^2 / c_\chi^2} - \frac{1}{2} \lambda_n q^2. \quad (28)$$

Expanding the square root for $|\dot{q}| \ll c_\chi$ gives

$$L_n[q] = \text{const} + \frac{1}{2} \dot{q}^2 + \frac{\kappa_n}{8c_\chi^2} \dot{q}^4 - \frac{1}{2} \lambda_n q^2 + \mathcal{O}(c_\chi^{-4}), \quad (29)$$

where

$$\kappa_n := \sum_v \psi_n(v)^4. \quad (30)$$

The corresponding Euler–Lagrange equation is

$$\left(1 + \frac{3\kappa_n}{2c_\chi^2} \dot{q}^2\right) \ddot{q} + \lambda_n q = 0. \quad (31)$$

Using the harmonic-balance ansatz

$$q(t) \approx a_n \cos(\omega_n t), \quad (32)$$

and keeping only the fundamental harmonic, one obtains the first amplitude-dependent correction

$$\omega_n^2(a_n) \approx \lambda_n \left(1 - \frac{3\kappa_n \lambda_n a_n^2}{8c_\chi^2} \right), \quad (33)$$

or equivalently

$$\omega_n(a_n) \approx \sqrt{\lambda_n} \left(1 - \frac{3\kappa_n \lambda_n a_n^2}{16c_\chi^2} \right). \quad (34)$$

The DBI correction is therefore a softening non-linearity: the frequency decreases as the mode amplitude grows.

In the canonical Q_8 basis, the quartic shape factor is explicit. For the one-dimensional nontrivial abelian sectors,

$$\kappa_{\rho_k} = \sum_{v \in Q_8} \left(\frac{1}{\sqrt{8}} \right)^4 = \frac{1}{8}, \quad k = 1, 2, 3. \quad (35)$$

For each canonical basis vector in the ρ_4 sector,

$$\kappa_{\rho_4} = 4 \left(\frac{1}{2} \right)^4 = \frac{1}{4}. \quad (36)$$

Hence the raw-amplitude correction is stronger in the non-abelian sector, reflecting the larger localisation of the corresponding basis vectors.

However, rewriting the correction in terms of the effective amplitude $A_n = a_n \|\psi_n\|_\infty$ gives

$$\omega_n^2(A_n) \approx \lambda_n \left(1 - \frac{3\lambda_n A_n^2}{8c_\chi^2} \frac{\kappa_n}{\|\psi_n\|_\infty^2} \right). \quad (37)$$

In the canonical Q_8 basis,

$$\frac{\kappa_{\rho_k}}{\|\psi_{\rho_k}\|_\infty^2} = \frac{(1/8)}{(1/\sqrt{8})^2} = 1, \quad \frac{\kappa_{\rho_4}}{\|\psi_{\rho_4}\|_\infty^2} = \frac{(1/4)}{(1/2)^2} = 1. \quad (38)$$

Therefore

$$\omega_n^2(A_n) \approx \lambda_n \left(1 - \frac{3\lambda_n A_n^2}{8c_\chi^2} \right) \quad (39)$$

in both the nontrivial abelian and non-abelian sectors. At fixed effective amplitude, the non-abelian sector thus remains less penalised by the weakly non-linear DBI correction, solely because $\lambda_{\rho_4} = 6$ is smaller than $\lambda_{\rho_k} = 8$.

4.6 Non-perturbative DBI envelope

The weakly non-linear reduction derived in Section 4.5 provides the leading amplitude correction to the dispersion relation. However, the discrete Dirac–Born–Infeld dynamics allows the single-mode problem to be solved exactly on the Q_8 Cayley graph.

Consider again a single eigenmode

$$\chi_v(t) = q(t)\psi_n(v),$$

with $\mathcal{L}\psi_n = \lambda_n\psi_n$ and $\|\psi_n\|_2 = 1$. Substituting this ansatz into the DBI Lagrangian Eq. (26) gives the exact one-mode Lagrangian

$$L_n[q] = -c_\chi^2 \sum_{v \in Q_8} \sqrt{1 - \frac{\psi_n(v)^2 \dot{q}^2}{c_\chi^2}} - \frac{1}{2} \lambda_n q^2. \quad (40)$$

Because the eigenmodes on the Q_8 Cayley graph have a simple support structure, the DBI sum can be evaluated explicitly.

Abelian sectors

For the nontrivial abelian representations ρ_1, ρ_2, ρ_3 , the normalized eigenmodes satisfy

$$|\psi_n(v)| = \frac{1}{\sqrt{8}} \quad \forall v \in Q_8.$$

The Lagrangian therefore becomes

$$L = -8c_\chi^2 \sqrt{1 - \frac{\dot{q}^2}{8c_\chi^2}} - \frac{1}{2}\lambda_n q^2. \quad (41)$$

Non-abelian sector

For the quaternionic sector ρ_4 in the canonical basis, the eigenmodes satisfy

$$|\psi_n(v)| = \begin{cases} \frac{1}{2} & \text{on four vertices,} \\ 0 & \text{on four vertices.} \end{cases}$$

The DBI sum therefore reduces to

$$L = -4c_\chi^2 \sqrt{1 - \frac{\dot{q}^2}{4c_\chi^2}} - 4c_\chi^2 - \frac{1}{2}\lambda_n q^2. \quad (42)$$

The constant term does not affect the dynamics, and the effective Lagrangian becomes

$$L = -4c_\chi^2 \sqrt{1 - \frac{\dot{q}^2}{4c_\chi^2}} - \frac{1}{2}\lambda_n q^2. \quad (43)$$

Energy conservation

Since the Lagrangian is time independent, the corresponding Hamiltonian is conserved. One obtains

$$E = 4c_\chi^2 \frac{1}{\sqrt{1 - \frac{\dot{q}^2}{4c_\chi^2}}} + \frac{1}{2}\lambda_n q^2. \quad (44)$$

At the turning points $\dot{q} = 0$, where $q = \pm a_n$, the energy becomes

$$E = 4c_\chi^2 + \frac{1}{2}\lambda_n a_n^2, \quad (45)$$

where a_n denotes the raw oscillation amplitude of the modal coordinate.

Velocity bound

The DBI square root imposes the kinematic bound

$$|\dot{q}| < 2c_\chi. \quad (46)$$

Expressed in terms of the field variable,

$$\chi_v = q\psi_n(v),$$

this corresponds exactly to the local flux constraint

$$|\partial_t \chi_v| \leq c_\chi.$$

Amplitude bound

The turning point occurs when the maximal vertex saturates the flux bound,

$$|\partial_t \chi_v| = |\dot{q}| |\psi_n(v)| = c_\chi.$$

Since $\|\psi_n\|_\infty$ is the maximal value of $|\psi_n(v)|$, this gives

$$a_n^{\max} = \frac{c_\chi}{\|\psi_n\|_\infty \sqrt{\lambda_n}}. \quad (47)$$

The corresponding effective amplitude

$$A_n = a_n \|\psi_n\|_\infty$$

therefore satisfies

$$A_n^{\max} = \frac{c_\chi}{\sqrt{\lambda_n}}. \quad (48)$$

Exact admissibility envelope

We thus obtain the non-perturbative result

$$A_n^{\max} = \frac{c_\chi}{\sqrt{\lambda_n}}, \quad (49)$$

which coincides exactly with the admissibility envelope derived from the linear bounded-flux argument.

The weakly non-linear correction therefore does not modify the location of the admissibility boundary. Instead, the DBI dynamics smoothly interpolates between the linear regime and a velocity-saturated regime as the amplitude approaches the bound.

This shows that the spectral admissibility envelope identified in Step 1 is not merely a linear approximation but the exact saturation boundary for the single-mode DBI dynamics on the Q_8 Cayley graph.

5 Main Result

Proposition 2 (Spectral admissibility hierarchy on Q_8). *In the undirected Cayley graph $\Gamma(Q_8, \{\pm i, \pm j, \pm k\})$, under the local bounded-flux constraint (8) and linearised dispersion, the maximal admissible effective amplitude satisfies*

$$A_{\rho_4}^{\max} = \frac{c_\chi}{\sqrt{6}} > \frac{c_\chi}{\sqrt{8}} = A_{\rho_k}^{\max}, \quad k = 1, 2, 3. \quad (50)$$

The non-abelian sector ρ_4 occupies a strictly larger admissible region in the space of effective mode amplitudes than any nontrivial abelian sector, by a factor

$$\frac{A_{\rho_4}^{\max}}{A_{\rho_k}^{\max}} = \sqrt{\frac{8}{6}} = \sqrt{\frac{4}{3}} \approx 1.155. \quad (51)$$

Proof. Direct substitution of $\lambda_{\rho_4} = 6$ and $\lambda_{\rho_k} = 8$ into the admissibility envelope (15). $\square$

Remark 4. *The trivial sector ρ_0 has $\lambda = 0$ and thus formally unbounded $A^{\max}$. It corresponds to the uniform zero mode of the connected graph. In the present analysis it is factored out, since we are interested only in non-trivial excitations above the connected background.*

6 Interpretation

Proposition 2, together with the basis-resolved Proposition 1, establishes that, at fixed substrate capacity c_χ , the non-abelian irreducible sector of the quaternion Cayley graph sustains a larger local effective amplitude before violating the bounded-flux constraint, even though its raw amplitude bound is lower than that of the nontrivial abelian sectors.

In the language of the Cosmochrony programme:

- The admissibility envelope $A_n^{\max} = c_\chi/\sqrt{\lambda_n}$ defines a *geometry of exclusion* in the space of relational modes: modes above the envelope are dynamically inadmissible.
- The non-abelian sector ρ_4 lies strictly below this envelope for a larger range of effective amplitudes than the nontrivial abelian sectors ρ_1, ρ_2, ρ_3 : the bounded-flux constraint is less restrictive for non-abelian modes, enlarging their admissible window by a factor of $\sqrt{4/3}$.
- This advantage is a statement about the physically relevant local variable $A_n = a_n \|\psi_n\|_\infty$, not about the raw amplitude a_n . In the canonical basis, the ρ_4 modes are more localised ($\|\psi\|_\infty = 1/2$) than the abelian ones ($\|\psi\|_\infty = 1/\sqrt{8}$), so the non-abelian sector is more constrained in raw amplitude even while remaining less constrained in effective amplitude.
- The weakly non-linear DBI correction is governed by the quartic shape factor $\kappa_n = \sum_v \psi_n(v)^4$. In raw amplitude, the canonical ρ_4 modes are more strongly penalised than the abelian ones because they are more localised. In effective amplitude, this localisation factor cancels on Q_8 , so the first DBI correction again favours the non-abelian sector through its smaller eigenvalue.
- This provides a minimal admissibility advantage for non-abelian relational modes, both at the linear level and at the first weakly non-linear DBI order, without any dynamical postulate beyond the local flux bound.

Remark 5 (Real field and complex eigenmodes). *The substrate field χ is taken real-valued. The natural eigenmodes of the 2D sector are matrix coefficients of ρ_4 , which are complex-valued. One may either work in $\ell^2(Q_8, \mathbb{C})$ and take real and imaginary parts to obtain real eigenmodes, or directly choose a real basis for each eigenspace. The local admissibility criterion is intrinsically defined, but its expression in terms of (a_n, ψ_n) depends on the chosen basis within a degenerate eigenspace through $\|\psi_n\|_\infty$. This basis dependence is resolved explicitly in Section 4.3.*

This result does not yet establish:

- that the non-abelian sector will dominate in a full dynamical simulation;
- the origin of the eigenvalue gap $\lambda_{\rho_4} = 6 < 8 = \lambda_{\rho_k}$ from the topological structure of Π alone;
- the fully non-linear Born–Infeld correction to the admissibility envelope near saturation beyond the weakly non-linear regime.

These remain open and constitute the next steps of the programme.

7 Synthesis: Spectral Filtering by Bounded Flux

The analysis performed in this note identifies the effective amplitude

$$A_n = a_n \|\psi_n\|_\infty$$

as the physically relevant variable controlling the saturation of the substrate.

7.1 Local Flux Constraint as a Spectral Filter

The bounded-flux condition

$$|\partial_t \chi_v| \leq c_\chi$$

acts locally on each vertex of the relational substrate.

For a single eigenmode

$$\chi_v(t) = a_n \psi_n(v) \cos(\omega_n t),$$

the constraint becomes

$$|\omega_n| a_n \|\psi_n\|_\infty \leq c_\chi.$$

Introducing the effective amplitude

$$A_n = a_n \|\psi_n\|_\infty$$

yields the admissibility envelope

$$A_n^{\max} = \frac{c_\chi}{\sqrt{\lambda_n}}.$$

A key consequence is that the localization of the mode does not alter the admissibility bound itself. Localization only determines how the raw amplitude a_n translates into a local flux impact on the substrate. The selection mechanism is therefore governed purely by the spectral eigenvalue λ_n .

7.2 DBI Softening as Dynamical Spectral Filtering

In the weakly non-linear regime the discrete Dirac–Born–Infeld dynamics produces an amplitude-dependent frequency

$$\omega_n^2(A_n) \approx \lambda_n \left(1 - \frac{3\lambda_n A_n^2}{8c_\chi^2} \right).$$

This simplified expression follows from the identity

$$\frac{\kappa_n}{\|\psi_n\|_\infty^2} = 1$$

which holds for all non-trivial sectors of the Q_8 Cayley graph.

The non-linear correction therefore depends only on the Laplacian eigenvalue. Modes with larger λ_n undergo stronger softening as the effective amplitude increases.

The relational substrate thus behaves as a dynamical low-pass spectral filter.

7.3 Conceptual Interpretation

The mechanism uncovered in this toy model can be summarized as

substrate capacity $\rightarrow$ spectral admissibility $\rightarrow$ nonlinear filtering $\rightarrow$ surviving sectors.

On the Q_8 Cayley graph the smallest non-trivial eigenvalue

$$\lambda_{\rho_4} = 6$$

belongs to the quaternionic non-abelian sector. This sector therefore possesses the largest admissible amplitude window and is the least affected by DBI softening.

In this sense, physical structures correspond to the spectral components that remain dynamically admissible under the bounded capacity of the relational substrate.

Matter may thus be interpreted as the dynamical residue of the spectral filtering imposed by the substrate.

The non-perturbative DBI reduction derived in Section 4.6 shows that the admissibility envelope

$$A_n^{\max} = \frac{c_\chi}{\sqrt{\lambda_n}}$$

is not merely a linear approximation. It coincides with the exact saturation boundary of the one-mode DBI dynamics on the Q_8 Cayley graph.

As the effective amplitude approaches this bound,

$$A_n \rightarrow A_n^{\max},$$

the modal dynamics enters a velocity-saturated regime in which the local flux constraint

$$|\partial_t \chi_v| \leq c_\chi$$

is approached at the maximally excited vertex.

In this regime the oscillation no longer grows in amplitude, and the dynamics becomes increasingly constrained by the substrate capacity. Energy injected beyond the admissible region cannot be sustained within the corresponding spectral component and is therefore redistributed toward modes with smaller spectral weight.

The bounded relational capacity of the substrate therefore acts not merely as a spectral cut-off but as a dynamical selector, favouring the most resilient spectral sectors.

7.4 Beyond the Q_8 Toy Model

The identity

$$\frac{\kappa_n}{\|\psi_n\|_\infty^2} = 1$$

is a special feature of the small and highly symmetric Q_8 graph.

On larger graphs this ratio will generally depend on the structure of the eigenmodes. The non-linear DBI correction may therefore amplify the spectral selection mechanism rather than acting uniformly across sectors.

This motivates extending the analysis to larger Cayley graphs, including Ramanujan families such as the Lubotzky–Phillips–Sarnak constructions on $\text{SL}(2, \mathbb{F}_p)$.

7.5 Next Step: Dynamical Survival

The next stage of the programme is to test this mechanism dynamically.

The proposed protocol is to initialise the substrate with a mixture of spectral modes close to the admissibility envelope and evolve the system under the discrete DBI dynamics.

If spectral cascades preferentially accumulate in the lowest admissible sector, this would demonstrate a genuine dynamical selection of symmetry sectors driven by the bounded relational capacity of the substrate.

Sections 2–7 established the spectral admissibility hierarchy on the Q_8 Cayley graph. We now identify the structural origin of this hierarchy and its possible connection with the $\text{SU}(2)$ geometry underlying spin.

8 Structural Origin of the Admissibility Hierarchy and the SU(2) Emergence Thread

8.1 The ratio $\sqrt{4/3}$ is not accidental

The admissibility hierarchy established in Proposition 2,

$$\frac{A_{\rho_4}^{\max}}{A_{\rho_k}^{\max}} = \sqrt{\frac{8}{6}} = \sqrt{\frac{4}{3}} \approx 1.155, \quad (52)$$

has a precise internal origin in the Q_8 Cayley spectrum. It originates entirely from the Laplacian eigenvalues $\lambda_{\rho_4} = 6$ and $\lambda_{\rho_k} = 8$ ($k = 1, 2, 3$). The ratio $8/6$ is in turn fixed by the adjacency eigenvalues $\mu_{\rho_k} = -2$ and $\mu_{\rho_4} = 0$ through $\lambda = d - \mu$ with $d = 6$.

The central observation is that $\mu_{\rho_4} = 0$: the quaternionic sector lies *exactly at the centre* of the adjacency spectrum $\mu \in [-d, d] = [-6, 6]$. This is a structural consequence of the character theory of Q_8 , not a numerical coincidence.

8.2 Why $\mu_{\rho_4} = 0$: character-theoretic argument

For any Cayley graph on a finite group G with symmetric generating set S , Schur's lemma gives the adjacency eigenvalue in the irreducible sector ρ as

$$\mu_\rho = \frac{1}{\dim \rho} \sum_{s \in S} \chi_\rho(s). \quad (53)$$

For the quaternionic representation ρ_4 of Q_8 , the generating set $S = \{\pm i, \pm j, \pm k\}$ (with $|S| = 6$, fixing the regularity degree $d = 6$) maps to the matrices

$$\rho_4(\pm i) = \pm \begin{pmatrix} i & 0 \\ 0 & -i \end{pmatrix}, \quad \rho_4(\pm j) = \pm \begin{pmatrix} 0 & 1 \\ -1 & 0 \end{pmatrix}, \quad \rho_4(\pm k) = \pm \begin{pmatrix} 0 & i \\ i & 0 \end{pmatrix}. \quad (54)$$

All six generator matrices are traceless: they are proportional to the Pauli matrices $\sigma_x, \sigma_y, \sigma_z$, which satisfy $\text{tr}(\sigma_\alpha) = 0$. Therefore

$$\chi_{\rho_4}(s) = \text{tr}(\rho_4(s)) = 0 \quad \forall s \in S, \quad (55)$$

and the character sum in Eq. (53) vanishes identically, giving $\mu_{\rho_4} = 0$.

This is a structural property: the generators of Q_8 map to traceless matrices under ρ_4 because Q_8 is generated by the unit quaternions i, j, k , whose matrix images are exactly the imaginary units of SU(2). The tracelessness is not a special feature of the graph but of the embedding $Q_8 \hookrightarrow \text{SU}(2)$.

The ratio $\sqrt{4/3}$ appearing in the admissibility hierarchy is not an accidental numerical feature of the toy model. It originates from the character structure of the quaternionic representation of Q_8 , for which the adjacency eigenvalue satisfies $\mu_{\rho_4} = 0$. The non-abelian sector occupies the central position of the adjacency spectrum, resulting in the smallest non-trivial Laplacian eigenvalue and hence the largest admissible amplitude window.

8.3 Q_8 as the minimal discrete encoding of spin

The connection between Q_8 and spin is not coincidental but topological. The rotation group $SO(3)$ has fundamental group $\pi_1(SO(3)) = \mathbb{Z}_2$, meaning that a rotation by 2π cannot be continuously contracted to the identity: one must rotate by 4π to return topologically to the starting point. This is the origin of half-integer spin.

The simply-connected double cover of $SO(3)$ is $SU(2)$, which is topologically the three-sphere S^3 :

$$SU(2) \cong \{a_0 I + ia_1 \sigma_x + ia_2 \sigma_y + ia_3 \sigma_z : a_0^2 + a_1^2 + a_2^2 + a_3^2 = 1\} \cong S^3. \quad (56)$$

The unit quaternions $\{q = a_0 + a_1 i + a_2 j + a_3 k : |q| = 1\}$ are in bijection with $SU(2)$ via the identification $i \leftrightarrow i\sigma_x$, $j \leftrightarrow i\sigma_y$, $k \leftrightarrow i\sigma_z$.

The group $Q_8 = \{\pm 1, \pm i, \pm j, \pm k\}$ is the smallest non-abelian subgroup of the unit quaternions, hence the smallest non-abelian subgroup of $SU(2)$. Its elements have the following interpretation as spin rotations: -1 is a 2π -rotation (the spinorial sign flip); $\pm i$, $\pm j$, $\pm k$ are π -rotations around the three coordinate axes. Q_8 therefore captures the minimal non-abelian subgroup structure compatible with the $SU(2)$ spin algebra: it contains both the 2π phase (-1) and the three independent rotation axes ($\pm i$, $\pm j$, $\pm k$).

The matrices $i\sigma_x$, $i\sigma_y$, and $i\sigma_z$ generate the Lie algebra $\mathfrak{su}(2)$ under commutation, so

$$Q_8 \xrightarrow{\rho_4} \text{Pauli algebra} \longrightarrow \mathfrak{su}(2) \longrightarrow SU(2) \longrightarrow \text{spin}. \quad (57)$$

The spectral result of this note—that bounded relational flux favours the ρ_4 sector—therefore says that bounded substrate capacity preferentially sustains precisely those modes encoding the spin structure of $SU(2)$.

8.4 Connection with the S^3 spectral invariant $\lambda_2/\lambda_1 = 8/3$

On the three-sphere S^3 , the Laplacian eigenvalues are $\lambda_\ell = \ell(\ell + 2)$ with multiplicity $(\ell + 1)^2$, giving

$$\ell = 0 : \lambda_0 = 0, \quad \ell = 1 : \lambda_1 = 3, \quad \ell = 2 : \lambda_2 = 8, \quad \ell = 3 : \lambda_3 = 15. \quad (58)$$

The fundamental spectral ratio of S^3 is

$$\left. \frac{\lambda_2}{\lambda_1} \right|_{S^3} = \frac{8}{3}. \quad (59)$$

This ratio is the primary spectral invariant used elsewhere in the Cosmochrony programme to characterise the S^3 geometry from the relational spectrum [?].

The ratio from the present Q_8 calculation,

$$\frac{\lambda_{\rho_k}}{\lambda_{\rho_4}} = \frac{8}{6} = \frac{4}{3}, \quad (60)$$

is related to the S^3 ratio by

$$\frac{\lambda_{\rho_k}}{\lambda_{\rho_4}} = \frac{1}{2} \cdot \left. \frac{\lambda_2}{\lambda_1} \right|_{S^3}. \quad (61)$$

This numerical relation suggests a possible connection between the coarse spectral structure of the Q_8 Cayley graph and the continuum spectrum of S^3 . At present this relation should be regarded as heuristic rather than as a derived theorem. The factor $1/2$ may reflect the extreme coarseness of the Q_8 approximation of S^3 , which has only 8 vertices. Testing whether this factor approaches unity for larger subgroups such as $2I$ is an open question.

Two independent computations in the Cosmochrony programme nonetheless point toward the same S^3 geometry: the spectral reconstruction ratio $8/3$ from [?], and the admissibility ratio $4/3$ derived here.

8.5 The SU(2) emergence programme

The Q_8 result motivates a systematic programme toward SU(2) emergence via spectral admissibility. The key chain of inclusions is

$$Q_8 \subset 2I \subset \text{SU}(2), \quad (62)$$

where $2I$ is the binary icosahedral group of order 120, the *largest* finite subgroup of SU(2), which is a double cover of $A_5 \cong \text{PSL}(2, \mathbb{F}_5)$.

The irreducible representations of $2I$ have dimensions $1, 2, 2', 3, 3', 4, 4', 5, 6$, which reproduce the low-spin content of SU(2): spin-1/2 (dim 2), spin-1 (dim 3), spin-3/2 (dim 4), spin-2 (dim 5), spin-5/2 (dim 6). The admissibility envelope $A_n^{\max} = c_\chi / \sqrt{\lambda_n}$ applied to Cayley graphs of $2I$ should therefore produce a *low-spin survival* hierarchy, in which lower-spin representations—carrying smaller Laplacian eigenvalues—are favoured.

Remark 6 (Heuristic status of low-spin survival). *The statement that lower-spin representations correspond to smaller Laplacian eigenvalues on $\text{Cay}(2I, S)$ is a structural hypothesis, not a theorem. It holds for Q_8 and is consistent with the general pattern that lower-dimensional representations tend to sit lower in the Cayley graph spectrum for groups approximating compact Lie groups, but it requires explicit verification on $2I$.*

The asymptotic step is provided by the Ramanujan graphs of Lubotzky, Phillips, and Sarnak [?], which are Cayley graphs of $\text{SL}(2, \mathbb{F}_p)$ designed to approximate S^3 with near-optimal spectral gap. These graphs have the property that all non-trivial eigenvalues satisfy $|\mu| \leq 2\sqrt{d-1}$ (the Alon–Boppana bound), making them optimal expanders. The bounded-flux admissibility mechanism studied here may therefore be connected to expander graph theory via the spectral gap.

The full programme is:

$$Q_8 \xrightarrow{\text{Step 4}} 2I \xrightarrow{\text{Step 5}} \text{SL}(2, \mathbb{F}_p) \xrightarrow{p \rightarrow \infty} \text{SU}(2). \quad (63)$$

At each stage the prediction is:

$$\text{bounded flux} \implies \text{spectral hierarchy} \implies \text{low-spin survival} \implies \text{SU}(2) \text{ emergence}. \quad (64)$$

8.6 Physical interpretation

The mechanism uncovered in this toy model has a precise physical reading within the Cosmochrony framework.

The non-abelian quaternionic sector has zero net correlation with the generators of the graph ($\mu_{\rho_4} = 0$): it neither oscillates *with* the graph structure nor against it. It is spectrally neutral with respect to the dynamics. This neutrality is the source of its resilience: the bounded-flux constraint imposes less penalty on modes that do not couple to the relational connectivity.

By contrast, the abelian sectors have $\mu_{\rho_k} = -2$, meaning that they are anti-correlated with the generators. They therefore pay a higher spectral price under bounded-flux filtering.

The mechanism is therefore not gauge symmetry, not dynamical postulate, and not fine-tuning. It is a purely structural consequence of the embedding $Q_8 \hookrightarrow \text{SU}(2)$ and the bounded capacity of the relational substrate:

$$\text{bounded relational capacity} \implies \text{non-abelian sector resilience}. \quad (65)$$

If this mechanism persists along the chain (62), the gauge symmetry group SU(2) would not need to be postulated. It would emerge as the limiting residue of spectral filtering under bounded substrate capacity.

Summary of Step 4. The Q_8 admissibility hierarchy is explained by the vanishing character sum $\sum_{s \in S} \chi_{\rho_4}(s) = 0$, itself a consequence of the tracelessness of the Pauli matrices. The resulting ratio $\sqrt{4/3}$ is numerically related by a factor $1/2$ to the fundamental spectral ratio $8/3$ of S^3 ; this relation is heuristic at present but both quantities point toward the same $SU(2) \cong S^3$ geometry. The Q_8 result is therefore not merely a toy model curiosity: it is the first discrete instance of a potential mechanism by which bounded relational capacity selects $SU(2)$ -type symmetry sectors.

9 Next Steps

The present note establishes the first four stages of the spectral admissibility programme on the Q_8 Cayley graph.

Step 1: spectral admissibility. The bounded-flux constraint $|\partial_t \chi_v| \leq c_\chi$ selects admissible modes through the effective amplitude $A_n = a_n \|\psi_n\|_\infty$, leading to the envelope

$$A_n^{\max} = \frac{c_\chi}{\sqrt{\lambda_n}}.$$

On Q_8 , this produces a strict hierarchy in which the quaternionic non-abelian sector ρ_4 admits the largest effective amplitude window.

Step 2: weakly non-linear DBI reduction. The discrete Dirac–Born–Infeld dynamics has been reduced to a single-mode effective equation, yielding the amplitude-dependent frequency

$$\omega_n^2(A_n) \approx \lambda_n \left(1 - \frac{3\lambda_n A_n^2}{8c_\chi^2} \right).$$

On the Q_8 Cayley graph the ratio $\kappa_n / \|\psi_n\|_\infty^2 = 1$ for all non-trivial sectors, implying that the non-linear softening depends only on the Laplacian eigenvalue. The spectral hierarchy obtained in Step 1 therefore remains stable under the leading DBI correction.

Step 3: non-perturbative DBI envelope. The next task is to move beyond the perturbative expansion in c_χ^{-2} and derive the full amplitude-dependent period from the exact one-mode Lagrangian (28). This can be obtained either through elliptic-integral methods or direct numerical quadrature. The resulting period–amplitude relation will determine the exact non-perturbative admissibility envelope replacing the linear boundary (15).

Step 4: structural origin of the admissibility hierarchy. Section 8 shows that the hierarchy observed on the Q_8 Cayley graph has a precise algebraic origin. The non-abelian quaternionic sector satisfies $\mu_{\rho_4} = 0$, placing it at the centre of the adjacency spectrum and yielding the smallest non-trivial Laplacian eigenvalue. This structural property explains why bounded relational flux favours the ρ_4 sector and suggests a possible connection with the $SU(2) \cong S^3$ geometry associated with spin.

Step 5: dynamical survival of spectral sectors. The next step is to simulate the full Q_8 substrate with the DBI constraint imposed locally at each vertex. The modal decomposition

$$A_n(t) = a_n(t) \|\psi_n\|_\infty$$

can then be tracked dynamically. The key question is whether spectral cascades concentrate energy in the sector with the largest admissible window, namely the non-abelian ρ_4 sector.

Step 6: extension to larger graphs. Finally, the toy model should be extended to Cayley graphs of larger groups approximating S^3 more faithfully. Natural candidates include the binary icosahedral group $2I$ (order 120) and families of Cayley graphs on $\mathrm{SL}(2, \mathbb{F}_p)$.

In particular, Ramanujan graph constructions such as the Lubotzky–Phillips–Sarnak graphs provide sequences with optimal spectral gap properties. Studying bounded-flux filtering on such graphs would test whether the spectral admissibility hierarchy persists in large expanders and may reveal a deeper connection between substrate capacity constraints and the spectral theory of expander graphs.